

# 1. Title : Natural Language Processing with Disaster Tweets

Twitter has become an important communication channel in times of emergency. The ubiquitousness of smartphones enables people to announce an emergency they're observing in real-time. Because of this, more agencies are interested in programatically monitoring Twitter (i.e. disaster relief organizations and news agencies).

But, it's not always clear whether a person's words are actually announcing a disaster.

The author explicitly uses the word "ABLAZE" but means it metaphorically. This is clear to a human right away, especially with the visual aid. But it's less clear to a machine.

In this competition, you're challenged to build a machine learning model that predicts which Tweets are about real disasters and which one's aren't. You'll have access to a dataset of 10,000 tweets that were hand classified. If this is your first time working on an NLP problem, we've created a quick tutorial to get you up and running.

## 2. Importing required Libraries

Using TensorFlow backend.

```
In [1]: import re
import numpy as np
import pandas as pd
from tqdm import tqdm
import matplotlib.pyplot as plt
import seaborn as sns
import nltk
from nltk.corpus import stopwords

import warnings
warnings.filterwarnings('ignore')

import keras
from keras.initializers import Constant
import tensorflow as tf
from tensorflow.keras.preprocessing.text import Tokenizer
from tensorflow.keras.preprocessing.sequence import pad_sequences
from tensorflow.keras import layers, Sequential
from tensorflow.keras import optimizers
from tensorflow.keras.optimizers import RMSprop
from tensorflow.keras.layers import LSTM
```

## 3. Loading the data

```
In [2]: train_data = pd.read_csv('input/train.csv', dtype={'id': np.int16, 'target': np.int16})
test_data = pd.read_csv('input/test.csv', dtype={'id': np.int16})
```

## 4. EDA

## 1) Check missing values

```
In [3]: train_data.isnull().sum()
```

```
Out[3]: id          0
keyword      61
location    2533
text         0
target       0
dtype: int64
```

Missing values exist in the keyword, location variables.

```
In [4]: miss_cols = ['keyword', 'location']


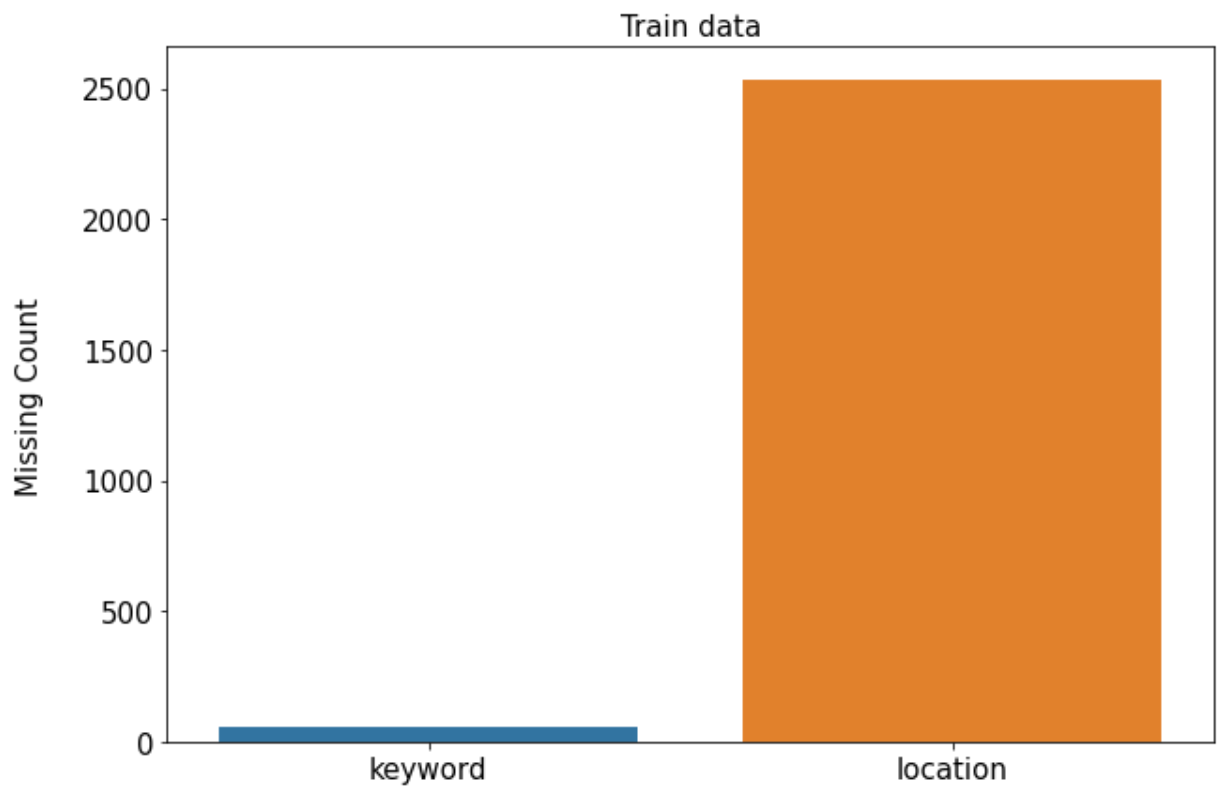
fig, axes = plt.subplots(2, figsize=(10, 15))

sns.barplot(x=train_data[miss_cols].isnull().sum().index, y=train_data[miss_cols].isr
sns.barplot(x=train_data[miss_cols].isnull().sum().index, y=train_data[miss_cols].isr

axes[0].set_ylabel('Missing Count', size=15, labelpad=20)
axes[0].tick_params(axis='x', labelsize=15)
axes[0].tick_params(axis='y', labelsize=15)
axes[1].set_ylabel('Missing Count', size=15, labelpad=20)
axes[1].tick_params(axis='x', labelsize=15)
axes[1].tick_params(axis='y', labelsize=15)

axes[0].set_title('Train data', fontsize=15)
axes[1].set_title('Test data', fontsize=15)

plt.show()
```



LOCATION variable has many missing values.

```
In [5]: train_data.groupby('target').count()['id']
```

```
Out[5]: target
0      4342
1      3271
Name: id, dtype: int64
```

There are more tweets with class 0 ( No disaster) than class 1 ( disaster tweets)

```
In [6]: train_x = train_data['text'].copy()
```

```
train_y = train_data['target'].copy()
```

0 is more than 1, but it is not much different.

## 2) Data Cleaning

Missing values must be processed before data analysis can be performed.

```
In [7]: stop = stopwords.words('english')
def clean(text):

    text = re.sub(r'httpWS+', ' ', text)

    text = re.sub(r'<.*?>', ' ', text)

    text = re.sub(r'#Ww+', ' ', text)

    text = re.sub(r'@Ww+', ' ', text)

    text = re.sub(r'Wd+', ' ', text)

    text = text.split()

    text = ' '.join([word for word in text if word not in stop])

    return text
```

### Removing Stop words

```
In [8]: train_x_cleaned = train_x.apply(clean)
train_x_cleaned.head()
```

```
Out[8]: 0          Our Deeds Reason May ALLAH Forgive us
1          Forest fire near La Ronge Sask. Canada
2    All residents asked 'shelter place' notified o...
3          , people receive evacuation orders California
4          Just got sent photo Ruby smoke pours school
Name: text, dtype: object
```

### Max Length

```
In [9]: max_len = max(train_x_cleaned.apply(len))
print('max length: {}'.format(max_len))
```

max length: 141

## 3) Tokenize

the text must be vectorized by generating a sequence of specified lengths for each tweet in the dataset. Using the Keras Tokenizer

```
In [10]: tokenizer = Tokenizer()
tokenizer.fit_on_texts(train_x_cleaned)
vocab_size = len(tokenizer.word_index) + 1
x = tokenizer.texts_to_sequences(train_x_cleaned)
x = pad_sequences(x, max_len, padding='post')
y = train_y
print('train_x_clean:', train_x_cleaned[4])
print('*'*50)
print('x:', x[5])
print('vocabulary size:{}'.format(vocab_size))
```

```

train_x_clean: Just got sent photo Ruby smoke pours school
*****
x: [ 318  45 1367  772 6083  478 1112  319   6   0   0   0   0   0
    0   0   0   0   0   0   0   0   0   0   0   0   0   0
    0   0   0   0   0   0   0   0   0   0   0   0   0   0
    0   0   0   0   0   0   0   0   0   0   0   0   0   0
    0   0   0   0   0   0   0   0   0   0   0   0   0   0
    0   0   0   0   0   0   0   0   0   0   0   0   0   0
    0   0   0   0   0   0   0   0   0   0   0   0   0   0
    0   0   0   0   0   0   0   0   0   0   0   0   0   0
    0]
vocabulary size:13947

```

```
In [11]: x.shape
```

```
Out[11]: (7613, 141)
```

```
In [12]: y.shape
```

```
Out[12]: (7613,)
```

## 5. Create Models

### model 1

GRU implementation with basic embedding layer

```

In [13]: epoch_size = 10
batch_size = 32
embedding_dim = 16
optimizer = optimizers.Adam(lr=3e-4)

model = Sequential([
    layers.Embedding(vocab_size, embedding_dim, input_length=max_len),
    layers.Bidirectional(layers.GRU(256, return_sequences=True)),
    layers.GlobalMaxPool1D(),
    # layers.Dense(128, activation='relu'),
    # layers.Dropout(0.4),
    layers.Dense(64, activation='relu'),
    layers.Dropout(0.4),
    layers.Dense(2, activation='sigmoid')
])
model.summary()

```

Model: "sequential"

Layer (type)	Output Shape	Param #
embedding (Embedding)	(None, 141, 16)	223152
bidirectional (Bidirectional)	(None, 141, 512)	420864
global_max_pooling1d (Global)	(None, 512)	0
dense (Dense)	(None, 64)	32832
dropout (Dropout)	(None, 64)	0
dense_1 (Dense)	(None, 2)	130

```
=====
Total params: 676,978
Trainable params: 676,978
Non-trainable params: 0
=====
```

## training the model

```
In [14]: model.compile(loss='sparse_categorical_crossentropy', optimizer = 'adam', metrics=['a
history = model.fit(x, y, epochs=epoch_size, validation_split=0.1)
```

```
Epoch 1/10
215/215 [=====] - 47s 208ms/step - loss: 0.5872 - accuracy:
0.6926 - val_loss: 0.5210 - val_accuracy: 0.7572
Epoch 2/10
215/215 [=====] - 46s 214ms/step - loss: 0.3567 - accuracy:
0.8581 - val_loss: 0.5055 - val_accuracy: 0.7598
Epoch 3/10
215/215 [=====] - 52s 243ms/step - loss: 0.2310 - accuracy:
0.9149 - val_loss: 0.5791 - val_accuracy: 0.7480
Epoch 4/10
215/215 [=====] - 60s 278ms/step - loss: 0.1572 - accuracy:
0.9466 - val_loss: 0.6171 - val_accuracy: 0.7441
Epoch 5/10
215/215 [=====] - 61s 285ms/step - loss: 0.1179 - accuracy:
0.9597 - val_loss: 0.7881 - val_accuracy: 0.7507
Epoch 6/10
215/215 [=====] - 62s 288ms/step - loss: 0.0967 - accuracy:
0.9666 - val_loss: 0.8090 - val_accuracy: 0.7231
Epoch 7/10
215/215 [=====] - 62s 287ms/step - loss: 0.0826 - accuracy:
0.9695 - val_loss: 0.8915 - val_accuracy: 0.7257
Epoch 8/10
215/215 [=====] - 62s 289ms/step - loss: 0.0719 - accuracy:
0.9712 - val_loss: 0.9324 - val_accuracy: 0.7257
Epoch 9/10
215/215 [=====] - 63s 293ms/step - loss: 0.0655 - accuracy:
0.9736 - val_loss: 1.0139 - val_accuracy: 0.7415
Epoch 10/10
215/215 [=====] - 64s 299ms/step - loss: 0.0594 - accuracy:
0.9747 - val_loss: 1.3628 - val_accuracy: 0.7323
```

```
In [15]: test_x = test_data['text'].copy()
test_x = test_x.apply(clean)
test_x = tokenizer.texts_to_sequences(test_x)
test_x = pad_sequences(test_x, max_len, padding='post')
```

## prediction

```
In [16]: test_pred = np.argmax(model.predict(test_x), axis=1)
print(test_pred)
```

```
[1 0 1 ... 1 1 1]
```

## Check the loss and accuracy

```
In [17]: history1 = history.history

trg_loss = history1['loss']
val_loss = history1['val_loss']

trg_acc = history1['accuracy']
val_acc = history1['val_accuracy']
```

```

epochs = range(1, len(trg_acc) + 1)

# plot losses and accuracies for training and validation
fig = plt.figure(figsize=(12,6))
ax = fig.add_subplot(1, 2, 1)
plt.plot(epochs, trg_loss, marker='o', label='Training Loss')
plt.plot(epochs, val_loss, marker='x', label='Validation Loss')
plt.title("Training / Validation Loss")
ax.set_ylabel("Loss")
ax.set_xlabel("Epochs")
plt.legend(loc='best')

ax = fig.add_subplot(1, 2, 2)
plt.plot(epochs, trg_acc, marker='o', label='Training Accuracy')
plt.plot(epochs, val_acc, marker='^', label='Validation Accuracy')
plt.title("Training / Validation Accuracy")
ax.set_ylabel("Accuracy")
ax.set_xlabel("Epochs")
plt.legend(loc='best')
plt.show()

```



In train data, it is a desirable state in which the loss is small and the accuracy value is high. However, in validation data, it is not convergent and the value is not good.

## model 2

Add Dropout layer

In [18]:

```

epoch_size = 10
batch_size = 32
embedding_dim = 16
optimizer = optimizers.Adam(lr=3e-4)

model2 = Sequential([
    layers.Embedding(vocab_size, embedding_dim, input_length=max_len),
    layers.Bidirectional(layers.GRU(256, return_sequences=True)),
    layers.GlobalMaxPool1D(),
    layers.Dense(64, activation='relu'),
    layers.Dropout(0.4),
    layers.Dense(2, activation='sigmoid')
])
model2.summary()

```

Model: "sequential\_1"

Layer (type)	Output Shape	Param #
embedding_1 (Embedding)	(None, 141, 16)	223152
bidirectional_1 (Bidirection	(None, 141, 512)	420864
global_max_pooling1d_1 (Glob	(None, 512)	0
dense_2 (Dense)	(None, 64)	32832
dropout_1 (Dropout)	(None, 64)	0
dense_3 (Dense)	(None, 2)	130

Total params: 676,978  
Trainable params: 676,978  
Non-trainable params: 0

```
In [19]: model2.compile(loss='sparse_categorical_crossentropy', optimizer = 'adam', metrics=['  
history = model2.fit(x, y, epochs=epoch_size, validation_split=0.1)
```

```
Epoch 1/10  
215/215 [=====] - 64s 289ms/step - loss: 0.5809 - accuracy:  
0.6887 - val_loss: 0.4593 - val_accuracy: 0.7874  
Epoch 2/10  
215/215 [=====] - 64s 299ms/step - loss: 0.3520 - accuracy:  
0.8622 - val_loss: 0.4848 - val_accuracy: 0.7979  
Epoch 3/10  
215/215 [=====] - 65s 303ms/step - loss: 0.2303 - accuracy:  
0.9174 - val_loss: 0.5995 - val_accuracy: 0.7572  
Epoch 4/10  
215/215 [=====] - 64s 298ms/step - loss: 0.1565 - accuracy:  
0.9470 - val_loss: 0.6014 - val_accuracy: 0.7887  
Epoch 5/10  
215/215 [=====] - 64s 298ms/step - loss: 0.1189 - accuracy:  
0.9599 - val_loss: 0.6471 - val_accuracy: 0.7730  
Epoch 6/10  
215/215 [=====] - 65s 302ms/step - loss: 0.0921 - accuracy:  
0.9677 - val_loss: 0.7340 - val_accuracy: 0.7480  
Epoch 7/10  
215/215 [=====] - 65s 303ms/step - loss: 0.0860 - accuracy:  
0.9704 - val_loss: 0.8525 - val_accuracy: 0.7546  
Epoch 8/10  
215/215 [=====] - 65s 302ms/step - loss: 0.0731 - accuracy:  
0.9714 - val_loss: 0.9728 - val_accuracy: 0.7533  
Epoch 9/10  
215/215 [=====] - 65s 303ms/step - loss: 0.0601 - accuracy:  
0.9752 - val_loss: 1.0421 - val_accuracy: 0.7572  
Epoch 10/10  
215/215 [=====] - 65s 301ms/step - loss: 0.0540 - accuracy:  
0.9759 - val_loss: 0.9407 - val_accuracy: 0.7559
```

```
In [20]: test_pred = np.argmax(model2.predict(test_x), axis=1)  
print(test_pred)
```

```
[1 0 1 ... 1 1 1]
```

```
In [21]: history2 = history.history  
  
trg_loss = history2['loss']  
val_loss = history2['val_loss']
```



```


trg_acc = history2['accuracy']
val_acc = history2['val_accuracy']

epochs = range(1, len(trg_acc) + 1)

# plot losses and accuracies for training and validation
fig = plt.figure(figsize=(12,6))
ax = fig.add_subplot(1, 2, 1)
plt.plot(epochs, trg_loss, marker='o', label='Training Loss')
plt.plot(epochs, val_loss, marker='x', label='Validation Loss')
plt.title("Training / Validation Loss")
ax.set_ylabel("Loss")
ax.set_xlabel("Epochs")
plt.legend(loc='best')

ax = fig.add_subplot(1, 2, 2)
plt.plot(epochs, trg_acc, marker='o', label='Training Accuracy')
plt.plot(epochs, val_acc, marker='^', label='Validation Accuracy')
plt.title("Training / Validation Accuracy")
ax.set_ylabel("Accuracy")
ax.set_xlabel("Epochs")
plt.legend(loc='best')
plt.show()

```



It is improved over model 1, but it is still not convergent and has a good value in validation data.

### model 3

LSTM implementation with basic embedding layer

```

In [22]: epoch_size = 10
batch_size = 32
embedding_dim = 16
optimizer = optimizers.Adam(lr=3e-4)

model3 = Sequential([
    layers.Embedding(vocab_size, embedding_dim, input_length=max_len, trainable=False),
    layers.SpatialDropout1D(0.2),

    layers.Bidirectional(layers.LSTM(64, recurrent_dropout=0.5, dropout=0.5, return_sequences=True)),
    layers.Bidirectional(layers.LSTM(64, recurrent_dropout=0.5, dropout=0.5)),
])

```

```

        layers.Dense(64, activation='relu'),
        layers.Dense(2, activation='sigmoid')
    ])
model3.summary()

```

Model: "sequential\_2"

Layer (type)	Output Shape	Param #
embedding_2 (Embedding)	(None, 141, 16)	223152
spatial_dropout1d (SpatialDr	(None, 141, 16)	0
bidirectional_2 (Bidirection	(None, 141, 128)	41472
bidirectional_3 (Bidirection	(None, 128)	98816
dense_4 (Dense)	(None, 64)	8256
dense_5 (Dense)	(None, 2)	130
Total params: 371,826		
Trainable params: 148,674		
Non-trainable params: 223,152		

```

In [23]: model3.compile(loss='sparse_categorical_crossentropy', optimizer = optimizer, metrics=
          history = model3.fit(x, y, epochs=epoch_size, validation_split=0.1)

```

```

Epoch 1/10
215/215 [=====] - 60s 250ms/step - loss: 0.6830 - accuracy:
0.5736 - val_loss: 0.6955 - val_accuracy: 0.5341
Epoch 2/10
215/215 [=====] - 54s 252ms/step - loss: 0.6789 - accuracy:
0.5744 - val_loss: 0.6871 - val_accuracy: 0.5341
Epoch 3/10
215/215 [=====] - 54s 252ms/step - loss: 0.6768 - accuracy:
0.5709 - val_loss: 0.6813 - val_accuracy: 0.5381
Epoch 4/10
215/215 [=====] - 54s 253ms/step - loss: 0.6747 - accuracy:
0.5771 - val_loss: 0.6806 - val_accuracy: 0.5315
Epoch 5/10
215/215 [=====] - 54s 253ms/step - loss: 0.6734 - accuracy:
0.5809 - val_loss: 0.6791 - val_accuracy: 0.5328
Epoch 6/10
215/215 [=====] - 55s 255ms/step - loss: 0.6722 - accuracy:
0.5780 - val_loss: 0.6776 - val_accuracy: 0.5459
Epoch 7/10
215/215 [=====] - 55s 256ms/step - loss: 0.6721 - accuracy:
0.5754 - val_loss: 0.6750 - val_accuracy: 0.5538
Epoch 8/10
215/215 [=====] - 55s 258ms/step - loss: 0.6698 - accuracy:
0.5849 - val_loss: 0.6762 - val_accuracy: 0.5486
Epoch 9/10
215/215 [=====] - 56s 260ms/step - loss: 0.6692 - accuracy:
0.5839 - val_loss: 0.6830 - val_accuracy: 0.5302
Epoch 10/10
215/215 [=====] - 57s 265ms/step - loss: 0.6692 - accuracy:
0.5812 - val_loss: 0.6751 - val_accuracy: 0.5525

```

```

In [24]: test_pred = np.argmax(model3.predict(test_x), axis=1)
          print(test_pred)

```

```
[0 0 0 ... 0 0 0]
```

```

In [25]: history3 = history.history

```

```

trg_loss = history3['loss']
val_loss = history3['val_loss']


trg_acc = history3['accuracy']
val_acc = history3['val_accuracy']

epochs = range(1, len(trg_acc) + 1)

# plot losses and accuracies for training and validation
fig = plt.figure(figsize=(12,6))
ax = fig.add_subplot(1, 2, 1)
plt.plot(epochs, trg_loss, marker='o', label='Training Loss')
plt.plot(epochs, val_loss, marker='x', label='Validation Loss')
plt.title("Training / Validation Loss")
ax.set_ylabel("Loss")
ax.set_xlabel("Epochs")
plt.legend(loc='best')

ax = fig.add_subplot(1, 2, 2)
plt.plot(epochs, trg_acc, marker='o', label='Training Accuracy')
plt.plot(epochs, val_acc, marker='^', label='Validation Accuracy')
plt.title("Training / Validation Accuracy")
ax.set_ylabel("Accuracy")
ax.set_xlabel("Epochs")
plt.legend(loc='best')
plt.show()

```



It is improved over model 2, but it is still not convergent and has a good value in validation data.

#### model 4

GRU implementation with basic embedding layer and 3 Dense layers

In [26]:

```

epoch_size = 10
batch_size = 32
embedding_dim = 16
optimizer = optimizers.Adam(lr=3e-4)

model4 = Sequential([
    layers.Embedding(vocab_size, embedding_dim, input_length=max_len),
    layers.Bidirectional(layers.GRU(256, return_sequences=True)),
    layers.GlobalMaxPool1D(),

```

```

layers.Dense(128, activation='relu'),
layers.Dropout(0.4),
layers.Dense(64, activation='relu'),
layers.Dropout(0.4),
layers.Dense(2, activation='sigmoid')
])
model4.summary()

```

Model: "sequential\_3"

Layer (type)	Output Shape	Param #
embedding_3 (Embedding)	(None, 141, 16)	223152
bidirectional_4 (Bidirection	(None, 141, 512)	420864
global_max_pooling1d_2 (Glob	(None, 512)	0
dense_6 (Dense)	(None, 128)	65664
dropout_2 (Dropout)	(None, 128)	0
dense_7 (Dense)	(None, 64)	8256
dropout_3 (Dropout)	(None, 64)	0
dense_8 (Dense)	(None, 2)	130

Total params: 718,066  
 Trainable params: 718,066  
 Non-trainable params: 0

In [27]:

```

model4.compile(loss='sparse_categorical_crossentropy', optimizer = optimizer, metrics=
history = model4.fit(x, y, epochs=epoch_size, validation_split=0.2)

```

```

Epoch 1/10
191/191 [=====] - 108s 558ms/step - loss: 0.6758 - accuracy:
0.5796 - val_loss: 0.6510 - val_accuracy: 0.5542
Epoch 2/10
191/191 [=====] - 108s 566ms/step - loss: 0.4853 - accuracy:
0.7851 - val_loss: 0.5271 - val_accuracy: 0.7485
Epoch 3/10
191/191 [=====] - 107s 563ms/step - loss: 0.3294 - accuracy:
0.8765 - val_loss: 0.5050 - val_accuracy: 0.7551
Epoch 4/10
191/191 [=====] - 107s 563ms/step - loss: 0.2313 - accuracy:
0.9243 - val_loss: 0.5755 - val_accuracy: 0.7571
Epoch 5/10
191/191 [=====] - 107s 560ms/step - loss: 0.1709 - accuracy:
0.9425 - val_loss: 0.6335 - val_accuracy: 0.7387
Epoch 6/10
191/191 [=====] - 107s 559ms/step - loss: 0.1401 - accuracy:
0.9548 - val_loss: 0.7627 - val_accuracy: 0.7249
Epoch 7/10
191/191 [=====] - 107s 559ms/step - loss: 0.8050 - accuracy:
0.7066 - val_loss: 0.6906 - val_accuracy: 0.5345
Epoch 8/10
191/191 [=====] - 108s 565ms/step - loss: 0.6802 - accuracy:
0.5757 - val_loss: 0.6886 - val_accuracy: 0.5345
Epoch 9/10
191/191 [=====] - 108s 566ms/step - loss: 0.6437 - accuracy:
0.6115 - val_loss: 0.6392 - val_accuracy: 0.6658
Epoch 10/10
191/191 [=====] - 107s 559ms/step - loss: 0.3609 - accuracy:
0.8578 - val_loss: 0.7054 - val_accuracy: 0.6875

```

```
In [28]: test_pred = np.argmax(model2.predict(test_x), axis=1)
print(test_pred)
```

```
[1 0 1 ... 1 1 1]
```

```
In [29]: history4 = history.history


trg_loss = history4['loss']
val_loss = history4['val_loss']

trg_acc = history4['accuracy']
val_acc = history4['val_accuracy']

epochs = range(1, len(trg_acc) + 1)

# plot losses and accuracies for training and validation
fig = plt.figure(figsize=(12,6))
ax = fig.add_subplot(1, 2, 1)
plt.plot(epochs, trg_loss, marker='o', label='Training Loss')
plt.plot(epochs, val_loss, marker='x', label='Validation Loss')
plt.title("Training / Validation Loss")
ax.set_ylabel("Loss")
ax.set_xlabel("Epochs")
plt.legend(loc='best')

ax = fig.add_subplot(1, 2, 2)
plt.plot(epochs, trg_acc, marker='o', label='Training Accuracy')
plt.plot(epochs, val_acc, marker='^', label='Validation Accuracy')
plt.title("Training / Validation Accuracy")
ax.set_ylabel("Accuracy")
ax.set_xlabel("Epochs")
plt.legend(loc='best')
plt.show()
```



The loss value is high and the accuracy is significantly different from the train result.

## model 5

GloVe embedded LSTM, RNN and add BatchNormalization

```
In [30]: embeddings_dictionary = dict()
embedding_dim = 100
```

```
glove_file = open('glove.6B/glove.6B.100d.txt', encoding='UTF8')
for line in glove_file:
    records = line.split()
    word = records[0]
    vector_dimensions = np.asarray(records[1:], dtype='float32')
    embeddings_dictionary[word] = vector_dimensions
glove_file.close()
```

```
In [31]: embedding_matrix = np.zeros((vocab_size, embedding_dim))
for word, index in tokenizer.word_index.items():
    embedding_vector = embeddings_dictionary.get(word)
    if embedding_vector is not None:
        embedding_matrix[index] = embedding_vector
```

```
In [32]: epoch_size = 10
batch_size = 32

model5 = Sequential([
    layers.Embedding(input_dim=embedding_matrix.shape[0],
                     output_dim=embedding_matrix.shape[1],
                     weights = [embedding_matrix]
    ),
    layers.Bidirectional(LSTM(64, return_sequences = True, recurrent_dropout=0.2)),
    layers.GlobalMaxPool1D(),
    layers.BatchNormalization(),
    layers.Dropout(0.5),
    layers.Dense(128, activation='relu'),
    layers.Dropout(0.5),
    layers.Dense(64, activation='relu'),
    layers.Dropout(0.5),
    layers.Dense(1, activation='sigmoid')
])
model5.summary()
```

Model: "sequential\_4"

Layer (type)	Output Shape	Param #
embedding_4 (Embedding)	(None, None, 100)	1394700
bidirectional_5 (Bidirection	(None, None, 128)	84480
global_max_pooling1d_3 (Glob	(None, 128)	0
batch_normalization (BatchNo	(None, 128)	512
dropout_4 (Dropout)	(None, 128)	0
dense_9 (Dense)	(None, 128)	16512
dropout_5 (Dropout)	(None, 128)	0
dense_10 (Dense)	(None, 64)	8256
dropout_6 (Dropout)	(None, 64)	0
dense_11 (Dense)	(None, 1)	65
Total params: 1,504,525		
Trainable params: 1,504,269		
Non-trainable params: 256		

```
In [33]: model5.compile(loss='binary_crossentropy', optimizer = RMSprop(learning_rate=0.0001),
history5 = model5.fit(x, y, epochs=epoch_size, validation_split=0.2)
```

```
Epoch 1/10
191/191 [=====] - 23s 105ms/step - loss: 0.9342 - accuracy:
0.5775 - val_loss: 0.6520 - val_accuracy: 0.6842
Epoch 2/10
191/191 [=====] - 21s 108ms/step - loss: 0.7763 - accuracy:
0.6123 - val_loss: 0.5710 - val_accuracy: 0.7452
Epoch 3/10
191/191 [=====] - 20s 107ms/step - loss: 0.7242 - accuracy:
0.6437 - val_loss: 0.5052 - val_accuracy: 0.7623
Epoch 4/10
191/191 [=====] - 21s 107ms/step - loss: 0.6672 - accuracy:
0.6726 - val_loss: 0.4816 - val_accuracy: 0.7800
Epoch 5/10
191/191 [=====] - 20s 107ms/step - loss: 0.6402 - accuracy:
0.6913 - val_loss: 0.4695 - val_accuracy: 0.7925
Epoch 6/10
191/191 [=====] - 21s 107ms/step - loss: 0.5941 - accuracy:
0.7158 - val_loss: 0.4607 - val_accuracy: 0.7938
Epoch 7/10
191/191 [=====] - 20s 106ms/step - loss: 0.5789 - accuracy:
0.7297 - val_loss: 0.4529 - val_accuracy: 0.7965
Epoch 8/10
191/191 [=====] - 20s 107ms/step - loss: 0.5639 - accuracy:
0.7363 - val_loss: 0.4493 - val_accuracy: 0.7965
Epoch 9/10
191/191 [=====] - 20s 107ms/step - loss: 0.5392 - accuracy:
0.7499 - val_loss: 0.4444 - val_accuracy: 0.8037
Epoch 10/10
191/191 [=====] - 20s 106ms/step - loss: 0.5385 - accuracy:
0.7493 - val_loss: 0.4398 - val_accuracy: 0.8037
```

```
In [34]: test_pred = np.argmax(model5.predict(test_x), axis=1)
print(test_pred)
```

```
[0 0 0 ... 0 0 0]
```

```
In [35]: history = history5.history

trg_loss = history['loss']
val_loss = history['val_loss']


trg_acc = history['accuracy']
val_acc = history['val_accuracy']

epochs = range(1, len(trg_acc) + 1)

# plot losses and accuracies for training and validation
fig = plt.figure(figsize=(12,6))
ax = fig.add_subplot(1, 2, 1)
plt.plot(epochs, trg_loss, marker='o', label='Training Loss')
plt.plot(epochs, val_loss, marker='x', label='Validation Loss')
plt.title("Training / Validation Loss")
ax.set_ylabel("Loss")
ax.set_xlabel("Epochs")
plt.legend(loc='best')

ax = fig.add_subplot(1, 2, 2)
plt.plot(epochs, trg_acc, marker='o', label='Training Accuracy')
plt.plot(epochs, val_acc, marker='^', label='Validation Accuracy')
plt.title("Training / Validation Accuracy")
ax.set_ylabel("Accuracy")
ax.set_xlabel("Epochs")
```

```
plt.legend(loc='best')
plt.show()
```



Both loss and accuracy values converge and the difference from the train data results tends to decrease. However, the loss value is still large and the accumulation value seems to need further improvement.

## model 6

GloVe embedded GRU

In [36]:

```
epoch_size = 10
batch_size = 32

model6 = Sequential([
    layers.Embedding(input_dim=embedding_matrix.shape[0],
                     output_dim=embedding_matrix.shape[1],
                     weights = [embedding_matrix]),
    layers.Bidirectional(layers.GRU(256, return_sequences=True)),
    layers.GlobalMaxPool1D(),
    layers.Dense(64, activation='relu'),
    layers.Dropout(0.4),
    layers.Dense(2, activation='sigmoid')
])
model6.summary()
```

Model: "sequential\_5"

Layer (type)	Output Shape	Param #
embedding_5 (Embedding)	(None, None, 100)	1394700
bidirectional_6 (Bidirection	(None, None, 512)	549888
global_max_pooling1d_4 (Glob	(None, 512)	0
dense_12 (Dense)	(None, 64)	32832
dropout_7 (Dropout)	(None, 64)	0
dense_13 (Dense)	(None, 2)	130



Total params: 1,977,550  
Trainable params: 1,977,550  
Non-trainable params: 0

---

```
In [37]: model6.compile(loss='sparse_categorical_crossentropy', optimizer = RMSprop(learning_rate=0.001),  
history6 = model6.fit(x, y, epochs=epoch_size, validation_split=0.1))
```


```
Epoch 1/10  
215/215 [=====] - 142s 651ms/step - loss: 0.5484 - accuracy:  
0.7254 - val_loss: 0.4592 - val_accuracy: 0.7940  
Epoch 2/10  
215/215 [=====] - 144s 667ms/step - loss: 0.4653 - accuracy:  
0.7942 - val_loss: 0.4424 - val_accuracy: 0.7966  
Epoch 3/10  
215/215 [=====] - 144s 671ms/step - loss: 0.4405 - accuracy:  
0.8098 - val_loss: 0.4555 - val_accuracy: 0.8005  
Epoch 4/10  
215/215 [=====] - 144s 669ms/step - loss: 0.4211 - accuracy:  
0.8165 - val_loss: 0.4485 - val_accuracy: 0.8045  
Epoch 5/10  
215/215 [=====] - 144s 670ms/step - loss: 0.4124 - accuracy:  
0.8219 - val_loss: 0.4305 - val_accuracy: 0.8097  
Epoch 6/10  
215/215 [=====] - 145s 677ms/step - loss: 0.4002 - accuracy:  
0.8323 - val_loss: 0.4333 - val_accuracy: 0.8150  
Epoch 7/10  
215/215 [=====] - 146s 677ms/step - loss: 0.3874 - accuracy:  
0.8384 - val_loss: 0.4287 - val_accuracy: 0.8163  
Epoch 8/10  
215/215 [=====] - 146s 679ms/step - loss: 0.3763 - accuracy:  
0.8415 - val_loss: 0.4253 - val_accuracy: 0.8176  
Epoch 9/10  
215/215 [=====] - 148s 687ms/step - loss: 0.3612 - accuracy:  
0.8489 - val_loss: 0.4335 - val_accuracy: 0.8071  
Epoch 10/10  
215/215 [=====] - 146s 681ms/step - loss: 0.3501 - accuracy:  
0.8548 - val_loss: 0.4339 - val_accuracy: 0.8071
```

```
In [38]: test_pred = np.argmax(model6.predict(test_x), axis=1)  
print(test_pred)
```

```
[1 0 1 ... 1 1 0]
```

```
In [39]: history = history6.history  
  
trg_loss = history['loss']  
val_loss = history['val_loss']  
  
trg_acc = history['accuracy']  
val_acc = history['val_accuracy']  
  
epochs = range(1, len(trg_acc) + 1)  
  
# plot losses and accuracies for training and validation  
fig = plt.figure(figsize=(12,6))  
ax = fig.add_subplot(1, 2, 1)  
plt.plot(epochs, trg_loss, marker='o', label='Training Loss')  
plt.plot(epochs, val_loss, marker='x', label='Validation Loss')  
plt.title("Training / Validation Loss")  
ax.set_ylabel("Loss")  
ax.set_xlabel("Epochs")  
plt.legend(loc='best')  
  
ax = fig.add_subplot(1, 2, 2)
```

```
plt.plot(epochs, trg_acc, marker='o', label='Training Accuracy')
plt.plot(epochs, val_acc, marker='^', label='Validation Accuracy')
plt.title("Training / Validation Accuracy")
ax.set_ylabel("Accuracy")
ax.set_xlabel("Epochs")
plt.legend(loc='best')
plt.show()
```



Both loss and accuracy values were improved in the desired direction.

## 6. Create a submission

Generate the results as a csv file.

```
In [40]: submission = pd.DataFrame({'id':test_data['id'], 'target':test_pred})
submission.to_csv('submission.csv', index=False)
```

In [ ]: